

UiO : Department of Informatics
University of Oslo

Group creation in a collaborative P2P channel allocation protocol

Identifying connected groups of access points

Hans Jørgen Furre Nygårdshaug

Master's Thesis, Spring 2018



Group creation in a collaborative P2P channel allocation protocol

Hans Jørgen Furre Nygårdshaug

11th February 2018

Abstract

Contents

1	Introduction	1
1.1	Motivation	2
1.2	Problem definition	3
1.3	Method	3
2	Background	4
2.1	Basic radio challenges	4
2.1.1	Collisions in wireless technology	4
2.2	MAC Layer	5
2.2.1	Carrier Sense	6
2.2.2	Collision Avoidance	6
2.2.3	Distributed Coordination Function	7
2.3	PHY Layer	7
2.3.1	PLCP Protocol Data Unit	8
2.3.2	Clear channel assessment	8
2.4	Radio Frequency Interference	9
2.4.1	Impact in 802.11	9
2.4.2	Countermeasures	9
2.5	Channels	10
2.6	Overview of Channel Allocation Algorithms	10
3	Related work	12
3.1	Brief survey of related work	12
3.1.1	Cisco RRM	12
3.1.2	DenseAP	12
3.1.3	HiveOS	13
3.1.4	ResFi	14
3.1.5	Distributed Clusteering	15
3.1.6	SCIFI	15

4 Data acquisition and structure	16
4.1 Motivation	16
4.2 Requirements	17
4.3 Program design	18
4.4 Data output and visual representation	19
4.5 WiGLE	20
4.5.1 REST API	22
4.5.2 The Haversine Formula	23
4.5.3 Data output	25
5 Access point clustering	26
5.1 Introduction	26
5.2 Problem overview	27
5.2.1 Centralized model	27
5.2.2 Distributed approach	28
5.3 Computation assumptions and requirements	29
5.4 Program design	29
5.4.1 Design choices	29
5.4.2 Group framework	30
5.4.3 Output file structure	30
5.5 Creating groups with clustering	31
5.5.1 Agglomerative Clustering	32
5.5.2 K-Nearest Neighbour Clustering	32
5.5.3 Results of on different topologies	35
5.5.4 Evaluation	35
5.6 Introducing splitting	36
5.6.1 Introduction to K-means	36
5.6.2 K-means splitting	37
5.6.3 Results	38
5.6.4 Evaluation	39
5.6.5 Minimum Cut	40
5.7 Evaluation	40
5.7.1 Simulation weaknesses	40
6 Group communication and state synchronization	42
6.1 Problem overview	42
6.2 Enabling technologies	43
6.2.1 Distributed consensus with Raft	43
6.2.2 Access point communication with ResFi	43
6.3 Protocol design	43
6.4 Protocol implementation	43

List of Figures

2.1	A and C sending to node B unknowingly at the same time, resulting in a collision	5
2.2	The timeline one frame transmission cycle in DCF mode	7
2.3	DSSS PHY PPDU format from IEEE Std 802.11-2016	8
2.4	Channel/frequency distribution	
	11	
4.1	Computing the interference between nodes	19
4.2	JSON output structure	20
4.3	Generated topology with random, uniform distribution and the interface for viewing topologies	21
4.4	Example of a Wigle API request	22
4.5	REST API response with AP data	23
4.6	Implementation of haversine distance	24
5.1	Group simulation file structure	31
5.2	Pseudocode sample of how the K-Nearest Neighbour Cluster- ing runs in a simulated environment	34
5.3	K-Nearest Neighbour Clustering on different topologies	35
5.4	K-means splitting on different topologies	38
5.5	Comparison with and without K-means splitting on Tynset . .	39
5.6	Comparison with and without K-means splitting on uniform distribution	40

Chapter 1

Introduction

The presence of Internet connected devices is still rapidly expanding. The latest forecast esitmates that there will be about 27 billion devices connected to the Internet by 2021 [6], while by the end of the 2017 the number connected devices in use will be 8.4 billion [9]. In 2020 households alone will be responsible for over 10 billion devices that are able to wirelessly connect to the home router [1], and wireless traffic will account for 63 percent of all IP traffic in the world [6]. Traditionally devices that are commonly connected to the Internet in a household are computers, phones, watches, smart TVs, audio systems, and lately also private network storage systems. As the era of Internet of Things (IoT) has rapidly descended upon us, increasingly also less obvious utilites are connected to the Internet. These devices may span everything from lights and HVAC systems to coffee machines, fridges and even toasters. Whereas in the early 2000s it was common to own one or two MAC-addresses per person, at the date of writing it is not unusual to possess a two digit number of MAC-addresses.

But it is not only the numbers of devices that has changed. The consumers' expectations and demands are also being altered over time. While ubiquitous connectivity is a buzzword often heard in the context of future 5G networks, ubiquitous access is already expected in modern households. The demand for Wi-Fi coverage extends through the entire home: the Internet radio in the garage, the video stream on the basement couch and the gaming laptop on the bedroom. All demands coverage and expects mobility at the same time. These devices also have something else in common that reflects the change that has happened the last years: demand for high data rate. In 2016 about 73 percent of all IP traffic orginated from video streams, and in 2021 this number is expected to increase all the way up to 82 percent [6].

This demand for continous streams of high bitrate video poses a challenge that can not only be solved by providing larger bandwidth. First of all, video

streams have high Quality of Service (QoS) demands. Buffering of videos and reduction of video quality negatively affects the subjective experience. Secondly, it is not unusual that different video flows occur simultaneously in a single residence. If such traffic is transmitted wirelessly it results in an inherently busy transmission medium. Because of the international spectrum allocation standards, Wi-Fi is largely restricted to a small number of channels. To overcome this challenge, the 802.11 protocol specifies a set of rules which allows only one device in the vicinity to transmit on a channel at the same time. This poses yet another challenge of distributing channels in a way that as few devices on the same channels as possible are close to one another.

In short, while the physical layer (PHY) traditionally has been upgraded to meet the new demands in bitrate (e.g. fiber to the home, and MIMO in 802.11ac) Wireless LAN connections struggle under the heavy impact of RF-interference which can not be solved by increasing the physical datarate capacity.

1.1 Motivation

Wireless LAN is deployed in almost all corporate buildings and residencies in the western world and increasingly also in the rest of the world. The use of these networks used to be limited to laptops that generated small amounts of data, but now the range of devices includes smart-phones, network storage devices, even in some cases servers. The deployment of Wi-Fi and the infrastructure has not changed much over the years to match the new demands and increased traffic, and in most places coverage is still the main concern, while QoS comes second.

Customers who are subscribed to high data rate service level agreements often find they can only receive a comparable data rate over wired LAN. When the Wi-Fi network in their home is constantly underperforming, it is not unusual a customer to upgrade the data rate of their agreement. If interfering networks were the perpatrators, usually increasing the bandwidth will have no effect. This leads to customers being largely frustrated with their Internet service providers, even though it is the underlying technology, and not the ISP which is at fault. A customer service representative or a tech-savvy consumer may manually switch the operating channel for a wireless access point. If this has no effect a customer might be encouraged to get a router which can transmit a more powerful signal. While this might give a short relief for the customer who was resourceful enough to deal with the problem, it in turn may trigger a chain reaction of even stronger interference

levels for the rest of the inhabitants in the surrounding area.

There has been done research on how centralized controllers can benefit the deployment of large enterprise networks ([15], [14] and [19]), but in this thesis we will address the emerging issue of deployment of Wi-Fi in residential areas and how routers and access points could organize themselves to allow distributed control and cooperation on channel allocation, where centralized controllers are impractical or impossible.

1.2 Problem definition

There are 3 non-overlapping channels on the 2.4GHz spectrum that 802.11 Wireless LANs uses. One of the common ways of selecting a channel is done by letting an access point sense which channel has the lowest interference levels, also called least congested channel search. When channels are selected in a selfish manner, where the only available information about the surrounding networks are obtained via subjective radio readings, it is highly unlikely that the distribution of channels in a confined area (e.g. an apartment block) becomes optimal. Ideally all the access points would be configured so that the channel distribution in this area lead to as few equal channels as possible being adjacent to one another. While this boils down to a NP-complete graph-coloring problem, it is not the focus of this thesis. The main challenge to be considered is the clustering problem of identifying access points that heavily impacts each other, and consider different approaches to how they can form groups in a distributed manner that seeks to intuitively identify nodes that lie within the same confined areas, as apartment buildings dense residential areas, etc.

1.3 Method

To be able to see if we can form groups in a way that creates clusters of nearby nodes we need to get some data to perform calculation on. We will both create synthetic data and use real world location data of access points to evaluate the performance of the group creations. Then we will take a look at possible solutions to let access points communicate with each other without any manual bootstrapping or previous association. Then we will consider the problems and challenges that have to be overcome in the process of creating and deploying an architecture as suggested in the thesis.

Chapter 2

Background

We will briefly introduce the relevant aspects of wireless technology and a selection of important concepts from the 802.11 standard, both on the link and physical layer.

2.1 Basic radio challenges

There are some challenges with wireless technologies that are harder to overcome than in wired transmission technologies like Ethernet.

The first step to achieve a successful transmission is making sure the receiving radio is within transmit range of the transmitter. The transmit range is decided by the power which the signal is transmitted at, the antenna gain, and the surrounding environment. If there are a lot of solid obstacles, like walls and ceilings, the signal is likely to have a very compromised range.

Even if the surrounding environment is mostly open space, the wireless signal becomes subject of attenuation, which is a physical property of electromagnetic waves that weakens the signal the longer it has travelled through a medium. When this medium is air, we refer to the phenomenon as Free Space Path Loss (FSPL). Attenuation limits the transmit range of a radio, and to transmit further it has to increase its transmission power.

2.1.1 Collisions in wireless technology

If there are other nodes nearby that are within radio B's sensing range that sends at the same time as A, radio B experiences radio frequency interference, and thus may not be able to correctly decode the signal of A. In 802.3 Ethernet this is graciously handled by collision detection in the CSMA/CD protocol. As each node can hear everyone on else on a cabled medium, and

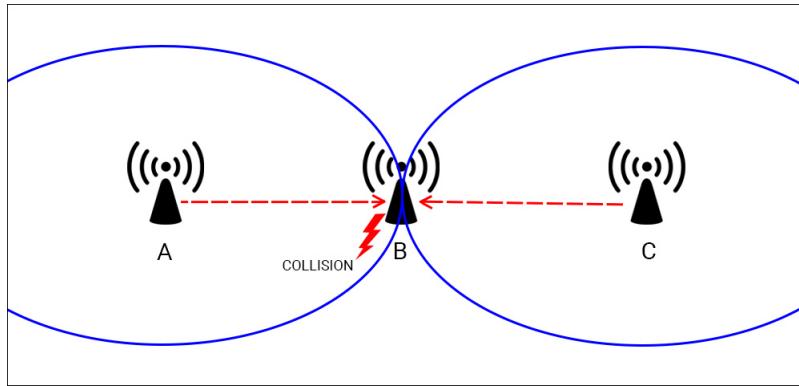


Figure 2.1: A and C sending to node B unknowingly at the same time, resulting in a collision

listening while transmitting is generally not a problem, a node can retransmit if a collision is detected. In wireless technologies it is not equally easy to listen while transmitting, and collision detection may be impossible because of the hidden terminal problem.

The hidden terminal problem

The hidden terminal problem is one of the major challenges for wireless technologies. When node A transmits a message to node B, it may not be able to sense what is going on on the opposite side of node B. If a node C transmits at the same time, this signal may not enter node A's sensing range, and hence go undetected by, even though a collision has happened near node B. This is illustrated in figure 2.1 where both A and C has unknowingly sent messages at the same time, and B has not been able to decode any messages because of the colliding messages.

2.2 MAC Layer

The 802.11 MAC layer implements the Carrier Sense Multiple Access / Collision Avoidance (CSMA/CA) protocol to control access to the medium. The CSMA/CA protocol is designed to operate in an entirely distributed fashion, where all stations connected to the same BSS operate on the same frequency without coordinated timeslots. As suggested by the name of the protocol, there are two basic operations in the CSMA/CA protocol: **Carrier Sensing (CS)** and **Collision Avoidance (CA)**.

2.2.1 Carrier Sense

In 802.11, carrier sensing (CS) is done in two ways

- **Physical carrier sensing** handled by the physical layer (PHY) as Clear Channel Assessment (CCA), which we will talk about in the PHY subsection.
- **Virtual carrier sensing** is a MAC layer mechanism in place to limit the number of times a node has to check the physical radio. When a valid 802.11 frame is decoded for a listening node, it can read the duration of the transmission from the MAC header. The frame with a duration is called a Network Allocation Vector (NAV). When a NAV is received the channel is marked as busy and the node will back off for the duration of the NAV.

2.2.2 Collision Avoidance

CSMA/CA attempts to avoid collisions in a network layout that includes hidden terminals. **Request To Send/Clear To Send** (RTS/CTS) is the function that allows CSMA/CA to some degree avoid the hidden terminal problem. By letting a node first ask the receiver if it is available for transmission (RTS), it prevents the node from sending the payload frame unless it receives Clear To Send (CTS) frame from the receiver first. The other mechanisms for collision avoidance are:

- **Interframe spacing** (IFS) is the amount of time the channel has to be idle before a sender can compete for channel access. To give priority to certain frame types, different types of frames can have different types of interframe spacing. The type of IFS is usually prefixed with the letter of the frame type. Organized by relative interval length, the different IFS are:
 - Short IFS (SIFS), before ACK, RTS, CTS.
 - DCF Mode IFS (PIFS), before RTS frames (or DCF data frames if RTS/CTS is disabled)
- **Exponential backoff** is what prevents two competing nodes from sending at the same time. When the channel is clear for DIFS time, a node has to wait another random number of milliseconds before transmitting. This is called backoff. The amount of backoff is randomly chosen from a contention window (CW). The contention window has a low start size, called CW_{min} . A node draws a backoff time in the range

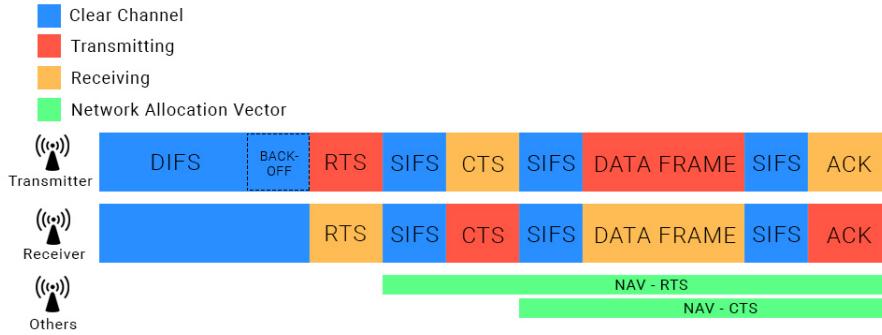


Figure 2.2: The timeline one frame transmission cycle in DCF mode

$(0, 2^n * CW_{\min})$, where n is the number of times the transmission has failed, beginning at $n = 0$, and $CW_{\min} < CW_{\max}$.

2.2.3 Distributed Coordination Function

The distributed coordination function is the main mode of operation in the 801.11 MAC layer, and is supposed to provide fair and reliable transmission for all nodes on the same network. Figure 2.2 shows the frame exchange that happens. The transmitter has to wait DIFS time, before drawing a number from the contention window and backing off that amount. As no other transmissions has begun during this time and the channel is still clear, the transmitter sends out an RTS frame. When received by the receiver it waits SIFS time before transmitting a CTS frame. The transmitter then sends his data frame, and waits for the ACK that indicates a successful transmission. The RTS/CTS mechanisms introduces extra overhead, and is sometimes turned off. The size of the payload and the number of stations on the network decides whether it is beneficial to have on or not [3].

2.3 PHY Layer

The physical layer in 801.11 is also divided in two sublayers. The upper sublayer is the Physical Layer Convergence Procedure (PLCP), responsible for CCA and acting as a common interface for MAC layer drivers. The lower sublayer is the Physical Medium Dependent (PMD) which is responsible for modulation and directly interfaces with the radio. It is responsible for transmitting the complete frames, and receive and decode incoming frames.

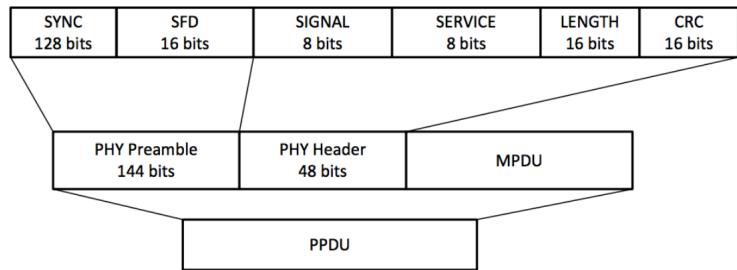


Figure 2.3: DSSS PHY PPDU format from IEEE Std 802.11-2016

2.3.1 PLCP Protocol Data Unit

The physical layer convergence procedure creates PLCP Protocol Data Units (PPDUs), that consists of 3 parts. The preamble, the header and the frame from the mac layer called MAC Service Data Unit (MSDU), see figure 2.3. The frame structure has a long and short format, and changes a little bit for High Rate DSSS (HR-DSSS), but contains mostly the same fields. As they are relevant for the thesis, we will briefly summarize the fields in the preamble and header:

- **Sync** bits are used to acquire the signal and synchronize timing.
- **SFD** stands for Start Frame Delimiter and is there to indicate the start of a frame.
- **Signal** the modulation type used to encode the MPDU and the data rate it is sent with.
- **Service** Reserved for future use.
- **Length** 16-bit fields that indicates the amount of time (in microseconds it will take to transmit the MPDU).
- **CRC** is the cyclic redundancy check that protects the fields signal, service and length.

2.3.2 Clear channel assessment

The PLCP layer also provides clear channel assessment. The purpose of clear channel assessment is to give information to the MAC layer if the carrier is available for transmission. There are primarily two ways the physical layer does CCA.

- CCA-ED (CCA-Energy Detect) detects signals that can not be decoded as a 801.11 frame, but is a disturbing signal on the channel. If the CCA-ED value exceeds a threshold, for instance 20 dB, then the CCA shall be indicated as busy by issuing a `PHY-CCA.indication(BUSY)` to the MAC layer.
- CCA-CS (CCA Carrier/Sense) detects a valid 801.11 frame and can properly decode the header fields of a valid PPDU frame. The channel gets marked as busy for as long as the length field in the PPDU header specifies, even if the observed signal is weaker than the ED threshold.

2.4 Radio Frequency Interference

Radio Frequency Interference (RF-interference) is the result of two or more signals being transmitted on the same frequency at the same time. A receiver will have problems deciding which parts of the signals belongs to which transmitter, and the signal may be altered to the extent that bits are changed or misrepresented. As the 2.4 GHz band that 802.11 utilizes is a part of the ISM band, channel noise or interference can come from sources such as microwaves, bluetooth devices or other WiFi entities.

2.4.1 Impact in 802.11

If the PPDU header gets corrupted by an interfering signal and can not be decoded, the PLCP layer issues a `PHY-RXEND.indication(CarrierLost)` to the MAC layer. According to CSMA/CA the station then has to wait EIFS (Extended Interframe Spacing) time before it can transmit a new frame. EIFS is defined as `ACK transmission time + SIFS + DIFS`. This is because the station that received the corrupt frame have no idea if any neighbouring station, received it correctly, and is about to transmit an ACK-frame. Additionally to waiting the minimum EIFS time, the station also has to wait for an idle channel indication from the PLCP layer again.

2.4.2 Countermeasures

Several countermeasures to limit the impact of RF-interference have been suggested. On the mac layer there is frame aggregation with individual headers. Frame aggregation in 802.11n is a technique that wraps several payloads under the same header and send them all when channel access is granted. This can improve the throughput if the channel is clear, but if the frame gets

corrupted during transmission an increased amount of data is lost. Therefore it can be beneficial to aggregate a frame with individual headers. Even though this gives a slightly increased processing and transmission overhead compared to regular aggregation where there is no individual headers, each frame can be selectively acknowledged. This means that only a few frames has to be retransmitted in case of corruption, and not the entire aggregation.

On the physical layer there are a couple of suggested countermeasures:

- Changing the transmission power levels
- Lowering data rates
- Adjusting CCA threshold
- Forward error correction
- Changing packet sizes

[[Fylle ut om de forskjellige]]

It is shown that the previous countermeasures only have limited impact, while changing the channel of an AP remains the most effective. [10].

2.5 Channels

802.11 b/g/n uses the range of frequencies from 2.400-2.500 GHz on the Industrial, Scientific and Medical (ISM) band. The increasingly popular 802.11n/ac also uses a range of 5 GHz frequencies on the Unlicensed National Information Infrastructure (UNII) band, which offers more frequencies [11]. Other than that the properties and challenges for the different bands are ultimately the same. The distribution of the frequencies on the 2.4 GHz band to the different channels is illustrated by figure 2.4. The frequencies listed are the center frequencies of each channel. In practice this means that an AP that transmits on one channel will interfere with close channels in both directions. Two channels are entirely non-interfering if they send on two frequencies f_1 and f_2 so that $|f_1 - f_2| > 0.025$. This means that there are in total three completely non-overlapping channels, 1, 6 and 11. The process of deciding which channel to transmit on is called channel allocation.

2.6 Overview of Channel Allocation Algorithms

CHAPTER 2. BACKGROUND

There exists many different channel allocation algorithms, and we'll introduce some relevant here.

Figure 2.4:
Channel/frequency distribution

CHNL_ID	Frequency (MHz)
1	2412
2	2417
3	2422
4	2427
5	2432
6	2437
7	2442
8	2447
9	2452
10	2457
11	2462
12	2467
13	2472
14	2484

Chapter 3

Related work

3.1 Brief survey of related work

3.1.1 Cisco RRM

CISCO offers a solution directed at enterprise networks [5], where implementing and managing a centralized controller is a more manageable task than in residential networks. Their architecture consists of one or more Wireless Lan Controllers (WLCs). The WLCs creates a group name which all APs part of the same RF-group is aware of. All APs broadcasts their RF-group name, along with the IP-address of their controller. Any other AP sharing the same RF-group name reports the incoming broadcast message to the controller, and then rebroadcasts the packet. This way the controller becomes aware of all APs that can hear each other, similar to the flooding routing mechanism, and they can form logical groups based on this information. If applicable the RF-groups perform leader election to decide which controller becomes the RF Group leader, but there can also a preconfigured leader. The RF group leader has the responsibility of running the relevant channel allocation algorithms and adjusting the radio power level of the APs.

3.1.2 DenseAP

DenseAP described in [14] aims to restructure the infrastructure of enterprise networks. The two fundamental changes they suggest is to deploy APs a lot denser (hence the name), and moving the task of associating clients with APs to a centralized controller. The reasoning behind the dense deployment of APs is that signals diminishes quickly in an indoor environment, and ideally an AP can always be associated with an AP in the close vicinity. The argument they use for moving the association decision away from the client

and over to a central controller, is that a client only uses signal strength as the metric to decide which to associate with. They emphasize that this can be suboptimal in conference and meeting room environments, where many clients seek to associate with an AP at the same time. If all clients pick the same AP it also means all clients will transmit on the same channel, and RF-interference can reduce the throughput on the medium. Their infrastructure consists of DenseAP access points (DAPs) and DenseAP Controllers (DCs). The DAPs send periodic reports to the DC, which contains information as RSSI measures, channel interference, and associated clients. Based on this information the DC decides which DAP each client should associate with, and also which channel each DAP should transmit on.

3.1.3 HiveOS

HiveOs [16], developed by Aerohive Networks offers distributed protocols and mechanisms to improve Wireless LANs in enterprise networks. The APs in the network are called HiveAPs, and they offer services such as

- Band steering: if a device can operate on the 5Ghz band, it will be forced to connect to the 5Ghz network to optimize the utilization of the radio spectrum.
- Load balancing: all HiveAPs have real-time information about how clients perform. If a client tries to associate with a new HiveAP, it will only be accepted if the new HiveAP has a low enough load to handle more clients. It would also know if other neighbouring APs are better suited to handle the load of the new client. This is achieved by withholding probe responses.
- Channel allocation: by using the Aerohive Channel Selection Protocol, HiveAPs tries to select the channel with the lowest co-channel interference. Their channel selection protocol uses 5 measurements, two static and 3 dynamic. The first static measurement is the number of nearby APs who are operating on the same channel. The more APs there are the higher the penalty, but the penalty per AP reduces as the number of APs increases. This is because the first one(s) are the most critical. The other static cost factor is what power level can be transmitted at the given channel, as this may differ on some 5GHz non-overlapping channels. The dynamic measurements are CRC-error rate, channel utilization, and the utilization of overlapping APs. All of these dynamic factors can penalize a channel with 0 to 3.5

3.1.4 ResFi

ResFi [24] is the only significant related work that directly aims to enable self-organized management in residential deployments of wireless LAN. Works such as [14], [5] and [16] are all directed toward enterprise networks. As we also aim to mainly concern ourselves with residential networks and their infrastructure, ResFi is especially interesting to look at. ResFi assumes that all access points have two interfaces, one connected to a wired backbone (e.g Internet), and another 802.11 compatible wireless interface. They also assume a Radio Resource Management Unit (RRMU), which in most residential homes is a router that controls the channel and power level of the antenna. In short ResFi enables connectivity between distributed access points without imposing a central controller on the access points, and without doing any modifications to hardware and drivers (like modifying standards or requiring proprietary equipment).

Operation

1. An AP that has just booted up beacons a frame on all available channels. This frame contains: a globally routable IP-address, port, and two cryptographic keys. One transient key for group communication, and the public key for the RRMU.
2. All APs that can hear the beacon responds with a probe response containing the same information as in 1.
3. When the scan is complete, the new AP can point-to-point communicate with all other APs using the wired backbone and the globally routable IP-address

Implementation

They implemented ResFi on Ubuntu 14.04. It adds vendor specific information elements to the MAC-header with the IP, port, and cryptographic data. This can be done Linux user space by using their modified version of hostapd [20]. As it runs on python, it could in practice be implemented on any Linux system. It uses a south-bound API to communicate with the RRMU, and a north-bound API to enable applications to use ResFi. ResFi itself provides no specific channel allocation mechanism, nor a group-creation/AP-clustering algorithm.

3.1.5 Distributed Clusteering

DCA and DMAC

3.1.6 SCIFI

SCIFI [2] is a centralized channel allocation protocol for infrastructure Wireless LANs that improves the traditional graph-coloring algorithm DSATUR. And while it shows that their central coordinator in fact improves the throughput compared to Wireless LANs that are not configured by SCIFI, the algorithm is just an algorithm for setting a channel in a preconfigured administrative domain. It does not deal with how the administrative domains (or collaborative groups as we call it) are defined.

Chapter 4

Data acquisition and structure

In this chapter we will consider how we can develop a tool that generates and visualizes the layout of network topologies based on constructed, artificial data and real, mined data. We will have great need for both artificial and real data in chapter 5 where we consider algorithms that enables distributed clustering of access points. This chapter also covers the storing of the data in question, how to structure it and visualize it both correctly and meaningfully in a web application.

4.1 Motivation

The data mining is important as we don't have the time nor the resources to organize a large enough testbed to purposefully evaluate the algorithms and suggestions that we will look at in the consecutive chapter. A sufficient testbed would require a large amount of routers (100-200 for a low scale test) in a small geographical area, preferably installed in residential apartment buildings. Not only would the creation of such a testbed require a lot of physical equipment, but there is also a lot of logical challenges that would have to be overcome, such as communication protocols and distributed consensus. These are problems we will address in chapter ???. While said challenges could be overcome, going directly from a concept idea to a large tesbed without having any empirical indication of which approaches have merit, is hardly advisable. Hence we gather data with the motivation of being able to perform realistic calculations and simulations, and with that enhance the believability of the results produced.

4.2 Requirements

SSID, current channel frequency, radio power, physical data rate and supported 802.11 standards are just a few properties of a single real-life access point. We are going to represent such access points in our network topology, but the access points we are representing will not be in the transmission (CSMA/CA) state where most of these properties are used. Our access points are in a state we can call the *group discovery state*, where the goal is to find or create a group to be a part of so that a channel can be assigned before transmission happens. Hence, to perform our computations, we don't have to consider many of these properties. Actually, we will only store each access point's SSID, because this is a practical way to uniquely identify them¹, and the list of other access points that can be seen through a WLAN scan with their observed signal strength. These are the only metrics we require, but unfortunately there is no publicly available data source that contains the subjective radio scans of a large amount of access points in the same area.

As a basis for our simulation we will represent nodes on a two-dimensional grid. Each dataset has to contain a set of nodes with two coordinates x and y so they can be positioned on the grid. When the nodes are placed on the grid they represent a network topology and it will be possible to compute an estimation of which nodes can hear each other on the radio, and add these to each node's SSID-scan list. This list contains the names of all the nodes it can hear, and how loud it is heard measured in $-dB_i$. Additionally the following parameters should be variable depending on each test scenario:

- Topology size with the possibility to give variable width of x- and y-axis as input arguments. These unit of the axes is meters.
- Number of nodes to place on the topology
- Minimum distance between nodes (in meters). This is only to avoid unlikely placement and extreme interference of nodes that are placed on top of each other.
- Minimum loudness measured in $-dB_i$ for a node to account another node as a neighbour (e.g -100 is too low for anyone to hear).

¹In the real world this is of course not the case, but we will enforce it in our computations. We could just as well call it a unique id.

4.3 Program design

The topology generation program consists of two main functionalitites.

The first functionality is creating a topology and generate nodes which are uniformly and randomly positioned on the network topology. The size of the topology, the number of nodes and the minimum distance between nodes are properties that are passed as input arguments to the program.

The second functionality is performing the calculation of which nodes can actually hear each other over the radio. We are assuming all nodes are transmitting with equal strength, and that the environment is flat and obstacle free. All the neighbouring nodes that can be heard by a node, is added to its list of neighbours, and it stores the $-dB_i$ value so it later can be shared with the group. The interference levels between access points are calculated by iterating through every access point. For each node N we record its x and y position, and then start a second iteration through the nodes. For each node n in the second iteration we calculate the distance d in meters between N and n using Euclidean distance. The formula for isotropic antennas is described by Friis [8], and can be used to derive the formula for free space path loss [23] that is as follows:

$$FSPL(dB) = 10 \log_{10} \left(\frac{(4\pi f d)^2}{c} \right)$$

Where $d = \text{distance}$, $f = \text{frequency}$ and $c = \text{constant}$. The constant c is used to account for different units. We will use the meters to denominate the distance, and megahertz for the frequency. The resulting formula which will be implemented in the program is

$$FSPL(dB) = 20 \log_{10} (f) + 20 \log_{10} (d) - 27.55$$

As we have to compute the distance from all nodes to every other node, the topology generation program has $O(n^2)$ complexity. A simplified version of the code can be seen in figure 4.1.

The resulting program, written in Python 3[7], contains an importable *topology class*. This way, for further testing we can use different data sources to get the positions of nodes, and only let the topology class compute the list of neighbours.

```

#In topology class
def measureInterference(self):
    for nodeSubject in self._nodes:
        for nodeObject in self._nodes:
            nodeSubject.calculateInterferenceTo(nodeObject)

#In node class
def calculateInterferenceTo(self, nodeObject):
    if self == nodeObject:
        return
    dist = round(self.distanceTo(nodeObject))

#If nodes have same coordinate, set high interference.
    if (dist == 0):
        dBi = -40
    else:
        dBi = self.measureDbi(dist) * -1

def measureDbi(self, dist):
    return (20 * math.log(self._frequency, 10)) +
(20 * math.log(dist, 10)) - 27.55

```

Figure 4.1: Computing the interference between nodes

4.4 Data output and visual representation

The result of the topology generation is stored in a JSON[12] data file. The data file contains the height and width of the the generated topology, as well as the number of nodes. A `nodes` object consists of as many JSON node-objects as there are nodes. Figure 4.2 illustrates the node structure and is an example of how a node with two neighbours will look.

```

1 {
2     "mapWidth": 100,
3     "mapHeight": 100,
4     "nodeCount": 3,
5     "nodes": [
6         {
7             "id": 1,
8             "posX": 50,
9             "posY": 50,
10            "ssid": "NODE1",
11            "neighbourCount": 2,
12            "neighbours": [
13                {
14                    "id": 0,
15                    "ssid": NODE2,
16                    "dbi": -77.23
17                },
18                {
19                    "id": 1,
20                    "ssid": NODE3,
21                    "dbi": -79.52
22                }
23            ],
24        }
25    ],
26    ...
27 }
28 ...
29 }
```

Figure 4.2: JSON output structure

The program can be run in the following way

```
./GenerateTopology.py -n 500 -w 100 -h 100 --space 10 --dbi 85
```

Which instructs the program to create a topology with 500 nodes. The topology should be 100 by 100 meters large, and there should be at least 10 meters between each node. The *dbi* parameter makes sure that only nodes which can be heard with a *-dB_i*-value of -85 or larger should appear in the neighbour list. The output from this specific run is a 23.2 MB large file containing the resulting topology-data in JSON.

By writing a simple HTML and JavaScript browser application, we can parse the JSON and visually represent the nodes on a grid. The result will look like what can be seen in figure 4.3

4.5 WiGLE

WiGLE (Wireless Geographical Logging Engine) [22] is a project started in 2001 which purpose is to gather information about wireless networks. All the information they collect is user submitted stored in their database. Anyone can download an Android app developed and published by WiGLE, and use



Figure 4.3: Generated topology with random, uniform distribution and the interface for viewing topologies

the app for wardriving,², then submit the data to WiGLE’s centralized database. The APs discovered can be viewed on an interactive map provided on WiGLE’s website. All the data can also be accessed through a public API which is especially interesting for us. Using their service is entirely free, but the amount of data that can be requested is throttled on a day-to-day basis. In the FAQ section on the website it’s written that the project openly supports research projects, so after contacting them they upgraded the account we will be using for the data requests to a higher daily data quota.

As the physical location of each access point is estimated using weighted-centroid trilateration, we can fetch the estimated position of each access point. However, the location data is not necessarily always accurate: if the measurements of an access point signal strength is done in multiple places (in a way that a line between the measurement locations surrounds the access point), the access point location becomes very precise. On the other hand, if the measurements are taken on only one side of the access point, the

²Wardriving is the act of tracking wireless networks using a laptop or a phone, and then store the information about each network.

```
https://api.wigle.net/api/v2/network/search?first=0&latrange1
=37.80846&latrange2=37.7467&longrange1=-122.5392&longrange2
=-122.3813&start=0
```

Figure 4.4: Example of a Wigle API request

measurements are heavily shifted towards that side. figure?

Even though the data is not perfect, it is more suited to impersonate the layout of real-life network topologies and give us a better understanding of how what access point clustering would look like in real-life.

4.5.1 REST API

WiGLE has a REST API that can be used to retrieve information from their database. The API responds to requests with JSON data. It gives access to a number of services such as user profile management, large-scale statistical information about the collected data and more specific network searches. An interactive guide to the public API is available at [21].

We are going to mine data for the topology generation program, hence the network search service is all we are interested in. There are several ways an access point can be retrieved using the network search. We can query specific SSIDs, a range of coordinates, only APs with a minimum signal level, or only query networks that are free to use. As we are only rebuilding the network topology to fit our tailored data structure, we will consider all access points and not filter on any parameters except for coordinates. By issuing an API a request for nodes between two latitudes and two longitudes, the response is an array with AP objects within that area. A request will look something like what can be seen in figure 4.4.

The parameters *latrange1* and *longrange1* are the coordinates that marks the beginning of the area we are interested in and *latrange2* and *longrange2* marks the end. We want information about the position of all access points within this range, however WiGLE returns at most 100 results per query. If there are more than 100 access points, we need to change the *start* parameter. This parameter tells WiGLE at which offset we want to begin fetching data from. For instance a start value of 0 means we fetch the first 100 access points, with indexes 0-99. A value of 100 means that we fetch the next 100 APs in the range 100 – 199 and so on. The JSON response for a successful request for one AP can be seen in figure 4.5.

The JSON-object in the response contains quite a lot of information about all of the APs retrieved. We wont be using most of this data, but it is worth noting that some of it could be used to perfect the search. For instance the

```

1 {
2 \subsection{Graph history}
3     "userfound": false,
4     "qos": 0,
5     "comment": null,
6     "lastupdt": "2015-12-22T17:49:34.000Z",
7     "bcninterval": 0,
8     "dhcp": "?",
9     "lasttime": "2015-12-22T17:49:15.000Z",
10    "trilong": 10.82792618,
11    "netid": "5C:9E:FF:2B:54:84",
12    "freenet": "?",
13    "trilat": 62.2816925,
14    "name": null,
15    "firsttime": "2015-12-22T20:55:01.000Z",
16    "type": "infra",
17    "ssid": "NETGEAR23",
18    "paynet": "?",
19    "wep": "2",
20    "transid": "20151222-00207",
21    "channel": 52
22 }

```

Figure 4.5: REST API response with AP data

last updated parameter could be a way to refine the access points retrieval, so that access points that has been long gone is not taken into consideration. The most valueable properties for us is *trilong* and *trilat*. As the names suggest these contain the position calculated by the weighted-centroid trilateration.

4.5.2 The Haversine Formula

As seen in the previous section, WiGLE provides data about the location of access points and a way for us to retrieve this data in a usable format. At this point we could create a program that operates directly on the longitudes and latitudes retrieved. This program would also need to be able to use the FPSL formula on global coordinates to compute the radio scans of each AP, and to properly visualize the data we would have to reinstall the nodes on a globe. This implementation requires a lot of extra labour, and seems especially redundant when we already have a working program that already has the aforementioned functionality. The problem is that our previous work relies on a two dimensional Cartesian coordinate system. To be able to reuse what we have have already built, the geographic coordinates has to be translated into coordinates in Euclidean space.

The haversine formula [18] is used to accurately compute the great-circle distance between two global coordinates.

$$d = 2R \sin^{-1} \left(\sqrt{\left(\sin\left(\frac{\varphi_2 - \varphi_1}{2}\right) \right)^2 + \cos(\varphi_1) * \cos(\varphi_2) * \sin\left(\left(\frac{\lambda_2 - \lambda_1}{2}\right)\right)^2} \right)$$

Where d is the distance between two latitudes φ_1 and φ_2 and two longitudes λ_1 and λ_2 . R is the radius of the sphere, which in this context is the earth's radius.

This can also be expressed with a two-parameter inverse tangent function [4], as long as neither of the parameters are zero.

$$\begin{aligned} a &= \left(\sin\left(\frac{\varphi_2 - \varphi_1}{2}\right) \right)^2 + \cos(\varphi_1) * \cos(\varphi_2) * \sin\left(\left(\frac{\lambda_2 - \lambda_1}{2}\right)\right)^2 \\ c &= 2 * \text{atan2}(\sqrt{a}, \sqrt{1 - a}) \\ d &= c * R \end{aligned}$$

The python implementation of the formula can be seen in figure 4.6. It is important to note that the degrees taken as input has to been converted to radians before inserting it in the formula.

```
def distanceBetweenCoordinates(lat1, lat2, long1, long2):
    deltaLon = math.radians(long2 - long1)
    deltaLat = math.radians(lat2 - lat1)
    a = (math.sin(deltaLat / 2))**2 + math.cos(math.radians(lat1)) * math.cos(
        math.radians(lat2)) * (math.sin(deltaLon/2))**2
    c = 2 * math.atan2(math.sqrt(a), math.sqrt(1 - a))
    R = 6371000
    d = round(c * R)
    return d
```

Figure 4.6: Implementation of haversine distance

This can be used to get the distance in meters between the boundaries of our area of interest. When the data is retrieved from WiGLE we can use the specified coordinates to compute the size of the cartesian coordinate system. The distance between the first latitude φ_{start} and the second latitude φ_{stop} is calculated. This is done by using the same longitude $\lambda_1 = \lambda_2$, and returns the size of the x-axis in meters. The opposite has to be done to get the size of the y-axis, using λ_{start} and λ_{stop} and equal longitudes. The same approach can be used to place each node on the coordinate system, where the first set of the coordinates is φ_{start} and λ_{start} , and the second set is the coordinates of the AP.

4.5.3 Data output

The resulting python program ??takes an output filename and two latitudes and two longitudes as input. It queries WiGLE’s API for all the APs between these coordinates. As long as the number of APs in the coordinate range does not exceed the daily data limit, all the data will be stored in a topology datastructure imported from the topology generation program in section 4.3. The APs will be placed correctly relative to each other, with real world distance between them, based on the coordinates of each AP.

wigle data figure, overlaid map

Chapter 5

Access point clustering

5.1 Introduction

While it at first sounds incredibly desirable to let the entire population of for instance New York's access points organize themselves in an optimal channel-plan, at second thought the idea may prove to be a little ambitious. Being limited by the NP-complete nature of DSATUR or similar graph coloring algorithms we have to set some reasonable constraints on the size of the *collaborating group*. The amount of nodes that will collaborate on the problem of finding an optimal channel distribution plan may not need to be very high.

Let us for a moment look at an ideal example: a city that only consists of small apartment buildings that are entirely isolated from interference produced in external networks. How could we build a group in such a case? Turns out it is not that hard. Each access point would be able to collaborate with every other observable access point. When computing channel distribution there would be no risk of surpassing the viable amount of nodes to compute the distribution for, as the group is limited to the apartment building. Sadly for us the world is not ideal, and it is not impossible that every access point in New York can observe eachother through a transitive relation. This means the size of the collaborating group would be vast and completely unviable.

This does not mean that creating reasonably sized groups of access points is impossible, but it poses a more difficult challenge. We need to filter out redundant nodes to create an approximated version of the ideal example.

Having a picture of the typical cityscape in mind, we know that buildings are naturally separated by streets, bridges, parks, and so on. Rural areas are similar, but networks are spaced more unevenly and usually only affecting each other in one plane.

Remembering attenuation, we know that RF-interference is also a property of distance. These pieces of information tells us that RSSI readings between access points in two separated buildings should be lower than readings between access points in adjacent apartments. In the following chapter we will consider how we can utilize the RSSI readings between access points to build a clustering algorithm that creates reasonable groups of access points.

5.2 Problem overview

The main problem we are dealing with is finding a way for access points to group together in clusters that are geographically close to each other. This can either be solved with a centralized coordinator or with a peer-to-peer distributed protocol. In this thesis we will be focusing on the distributed approach, but first we will take a look at the pros and cons of each approach.

5.2.1 Centralized model

Let us briefly envision how a centralized model might look. As we want to propose a solution that works across private consumer networks as well as larger corporate networks, we can not assume that all networks are under the same administrative domain. This is where the centralized approach is limited, and why most solutions presented in the related work section is usually only deployed under the same extended service set (ESS).

Just like a distributed version, a centralized controller is also restricted by the NP-complete nature of all graph colouring algorithms. To be able to compute a channel plan using heuristics, it has to identify groups of nodes that impacts each other severely.

The advantage of the centralized model is the ability to have the full picture of access points readily available. Let us assume the controller is placed at an ISP or a router manufacturer. The controller could potentially know exactly where in the world the access points are placed. Creating groups would be as simple as dividing within naturally separative geographical barriers like roads, streets and buildings.

For a central controller to be able to identify APs that are near each other and compute the optimal channel distribution, APs need to report their radio readings to the controller. The controller then needs to identify which of all the observed APs it can control, and which is not controllable and has to be treated as noise. There would be no requirement for communication in between nodes, which reduces complexity a lot.

A major drawback of the centralized approach is that nodes that are near each other have to be under the same controller. If there are too many nodes around that can not be controlled by the controller, all nearby access points would be treated as noise and regular channel allocation schemes would have to be applied. Even though in Norway there is a tendency for apartments in the same building to have the same ISP, there is no guarantee for that to always be the case.

5.2.2 Distributed approach

We will continue considering a distributed approach as the main topic of the thesis. In reality this idea has some major challenges that we have to overcome. Here are the most prominent ones:

1. The group creation itself. Based on the RSSI-readings and the communication channel, nodes have to be able to organize themselves in a tight cluster of nodes that includes all nodes that impacts each other most severely. This problem essentially boils down to a clustering problem. This is the main issue we will try to find a reasonable solution for.
2. The communication channel. Nodes have to be able to communicate with each other. The 802.11 standard describes no protocol for communication between access points that are not on the same extended service set. This communication would have to happen on the network layer, preferably over TCP. This communication channel would have to convey messages that enables group creation, most likely a custom protocol.
3. State synchronization. As all nodes have to compute a channel plan for the group, they all need to have synchronized information about the members of the groups and possibly also every other node's RSSI-readings.
4. Channel plan computation. Even though the centralized approach would have a similar problem, it will be more difficult to solve in a distributed fashion as it is reasonable to assume that an access point has limited computational power. This point extends past the scope of the thesis, but algorithms like DSATUR (insert ref) and SCIFI are examples of algorithms that computes channel plans as a graph colouring problems.

5.3 Computation assumptions and requirements

Moving forward in the next chapters we will be looking at possible ways for nodes to organize themselves into groups. We will be using the data we fetched in chapter 4. The data provides a topology of nodes where all nodes have a list of neighbouring nodes. This neighbour list contains the appropriate computed RSSI-readings for each node. Based on the RSSI-values of their neighbours, the nodes should be able to create groups that resembles clusters. It is important that the border between two different groups are placed in such a way that interference between the two groups are as minimal as possible. To simplify the computation procedure and to reduce the workload for building the simulation program we are going to make a number of assumptions.

1. Assumption 1: All nodes involved in the simulation also runs the group creation algorithm.
2. Assumption 2: When a node is observed over radio, it is also known how to directly contact the node (e.g. via TCP). This assumption also implies that there is implemented a protocol that lets nodes communicate and exchange information about their RSSI-readings and group membership. For now we will assume this is true, but in chapter (ref) we will address the issue.
3. Assumption 3: All nodes in a group are completely synchronous, and always have an equal image of the state of the group at a given time. Sharing information with the group is instant. This is also an issue that will be addressed in chapter 3.

5.4 Program design

5.4.1 Design choices

The program is designed to be modified so it can accomodate different algorithm types without changing the fundamental framework. Everything regarding group computation is implemented in Python 3 [7], and is designed to parse the output from the data generation program and use it directly. Group computation should be a fast operation, while the data generation (or fetching) is a slow process and should only be done once every time a new data set is required. The program is not parallelized to keep results consistent and to easier debug and locate program errors.

5.4.2 Group framework

The group framework consists of 3 classes with different responsibilities:

- **Group.** An object of this class is an abstraction of all the nodes that belong to the same group. Because we assumed that all nodes have equal information about group membership and RSSI-readings, we can store all the members of each group in a list and let the group object act as a unified entity on behalf of the entire group. All the interfaces and logic for forming groups is placed within this class. A method named `iteration` has the responsibility of triggering the appropriate action based on the state of the group. For instance adding nodes, removing nodes or merging the group with another. This is the part of the code that has to be changed when implementing and changing algorithms. If an action was performed the method returns 1, else it returns 0.
- **GroupCollection.** An object of this class contains all the groups used in a simulation. Its main functional responsibility is looping through all the groups and calling the `iteration` method of each group once. This is done in the GroupCollection's `iterate` method. It accumulates the amount of changes done in all the groups, by adding the return values of the Group object's `iteration` method. It also handles the destruction of groups, and bootstrapping of newly created groups.
- **Simulation.** An object of this class handles the bootstrapping of groups, where all nodes (given in the input file) are parsed and put in their own grown group. Consequently, at the beginning of each computation the amount of groups is equal to the amount of nodes. The Simulation class is also responsible for starting and stopping the simulation itself. After bootstrapping all the groups, the Simulation enters a loop where it calls the `iterate` method in the GroupCollections object once every run. All the groups have converged and reached a steady state once the amount of changes returned by the GroupCollection's `iterate` method is equal to zero. The results are written to file.

flowchart

5.4.3 Output file structure

The results of the group computations are written to file. The results does not only contain the resulting division of groups, but to be able to recreate the simulation visually, the results contains the topology of all nodes and their group membership for each iteration in the simulation process. This

means that we can step-by-step recreate the simulation. The data is stored as JSON, and the structure can been seen in figure 5.1. Having the data stored as JSON means that the data is language independent. This allows us to either implement a parser in python for the data and use `matplotlib` to visualize it, or we can use another applicaton to visualize the data. Since we already have the topology visualizer written in HTML and JavaScript from chapter 4, we can extend this program to let us upload a group creation output file.

```
1 "iterations": {
2     "0": {
3         "0": {
4             "groupName": "GROUP 0",
5             "members": {
6                 "0": "NODE 0"
7             },
8             "memberCount": 1
9         },
10        "1": {
11            "groupName": "GROUP 1",
12            "members": {
13                "0": "NODE 1"
14            },
15            "memberCount": 1
16        }
17    },
18    "1": {
19        "0": {
20            "groupName": "GROUP 1",
21            "members": {
22                "0": "NODE 0",
23                "1": "NODE 1"
24            },
25            "memberCount": 2
26        }
27    }
28 }
```

Figure 5.1: Group simulation file structure

This particular simulation had two iterations. In the first iteration there were two groups, each with one member node. In the second iteration the two groups has merged to one group, now containing both nodes.

5.5 Creating groups with clustering

In this section we will consider and evaluate a way to divide access points in groups.

5.5.1 Agglomerative Clustering

Description

Agglomerative clustering [13] is a variant of hierarchical clustering where clusters are pairwise greedily merged. The starting number of clusters is the same as the number of points in the data set. Based on a specified distance metric, the two clusters that are closest to one another is merged into combined cluster. The clustering is complete when there is only one cluster left. It falls under the category of hierarchical clustering, and a dendrogram is iteratively being built for each merge, containing all past merges of the clusters. An example of agglomerative clustering can be seen in figure??.

Evaluation

Agglomerative clustering provides us with the basic idea for a an algorithm we can use to create access point groups. However, there are two points where agglomerative clustering fails to meet our requirements:

- There is no upper limit to the amount of nodes in each cluster
- As the clustering algorithm will be running distributed on each cluster, a decision has to be made with only the local distance observations at hand. Assuming that there exists two groups that finds each other the closest (most disturbing), might prove a fallacy. RSSI is observed locally, and while one node observes another with a given signal strength, the mutual observation might be a different one. This could create a deadlock where there exists no combination of group pairs where both observes the other as the closest group.

The first point has the trivial solution of checking the number of nodes before accepting a merge, and would only be a slight implementation change of agglomerative clustering. The second point forces us to change the paradigm of always finding two mutually closest groups.

5.5.2 K-Nearest Neighbour Clustering

Description

To resolve the problems discussed in the previous section, we will suggest a slightly different algorithm and refer to it as K-Nearest Neighbour Clustering. In this method, similar to agglomerative clustering, in the beginning there are equally many clusters as there are nodes. Instead of looking for pairs that are mutually close, each cluster seeks to merge with the cluster that

is closest. The distance is defined by the distance metric. The merge will always happen as long as the resulting cluster does not contain a higher node count than K.

Distance metric

The distance metric lets us specify exactly what the distance between two clusters signifies. As each cluster usually consists of several nodes, there are multiple options for what the distance between two clusters could be defined as. The distance metrics used in hierarchical clustering are typically either

- The average distance between the nodes of each cluster
- The minimum distance between the nodes of each cluster
- The maximum distance between the nodes of each cluster

The metric we will use is the -dBi value of neighbouring nodes. Every node in a cluster tracks the observed -dBi value of all other nodes that are not a member of the cluster. The cluster containing the node with the highest -dBi value will be merged together with original one. This is similar to the minimum distance metric.

Implementation

Each node begins by identifying itself as a member of a group that only contains itself. Let us call this group a . Group a loops through the radio readings of every member of the group, and picks the node with the highest observed -dBi value to contact. In the beginning there is only one node in group a . Hence in the first iteration this node's radio readings alone will decide which group to merge with. The neighbour node, which we will call B , is the node that disturbs group a the most. B is a member of group b . In other words, group a wants to merge with group b to create a larger group that contains node B .

A merge happens in the following way: the members of the two groups exchange information about all their member nodes and their radio readings and combine the information. As the data is now identical for all the members of both groups and they can make identical choices it means that they are part of the same group.

In our simulation this is as easy as combining two Group objects into one. A pseudocode sample of an implementation can be seen in 5.2. In a real world implementation there would have to be a supporting protocol to

enable the flow of information and synchronization of data. These specific challenges will be addressed in chapter insert chapter.

```

allGroups = []
K = 120;

for node in topology: //Initialize groups
    g = new Group()
    g.members.append(node)
    allGroups.append(g)

while True: //Run as long as there are changes
    changes = 0;
    for group in allGroups:
        globalMaxDbi = -INFINITE
        globalClosestNode = None

        for node in group:
            localMaxDbi = -INFINITE
            localClosestNode = None

            for neighbour in node.neighbours:
                if neighbour.dbi > localMaxDbi:
                    localMaxDbi = neighbour.dbi
                    localClosestNode = neighbour

            if (localMaxDbi > globalMaxDbi):
                globalMaxDbi = localMaxDbi
                globalclosestnode = localclosestnode

    var groupToMergeWith = globalclosestnode.group
    if (length(group.members) + length(groupToMergeWith.members)) < K:
        var newGroup = new Group()
        newgroup.members = group.members + mergeWith.members
        allGroups.append(newgroup)
        allGroups.remove(mergeWith)
        allGroups.remove(group)
        changes = 1
    if (changes == 0):
        break

```

Figure 5.2: Pseudocode sample of how the K-Nearest Neighbour Clustering runs in a simulated environment

We can not always accept merges, else we would end up with a group that spanned the entire topology. That is why we define a maximum threshold for the amount of members a group can have, referred to as K. If the sum of members in two groups that wants to merge exceeds K, the merge is aborted and no changes is reported to have happened for either group. This means that the simulation algorithm converges when no groups remain that are small enough to merge with another.

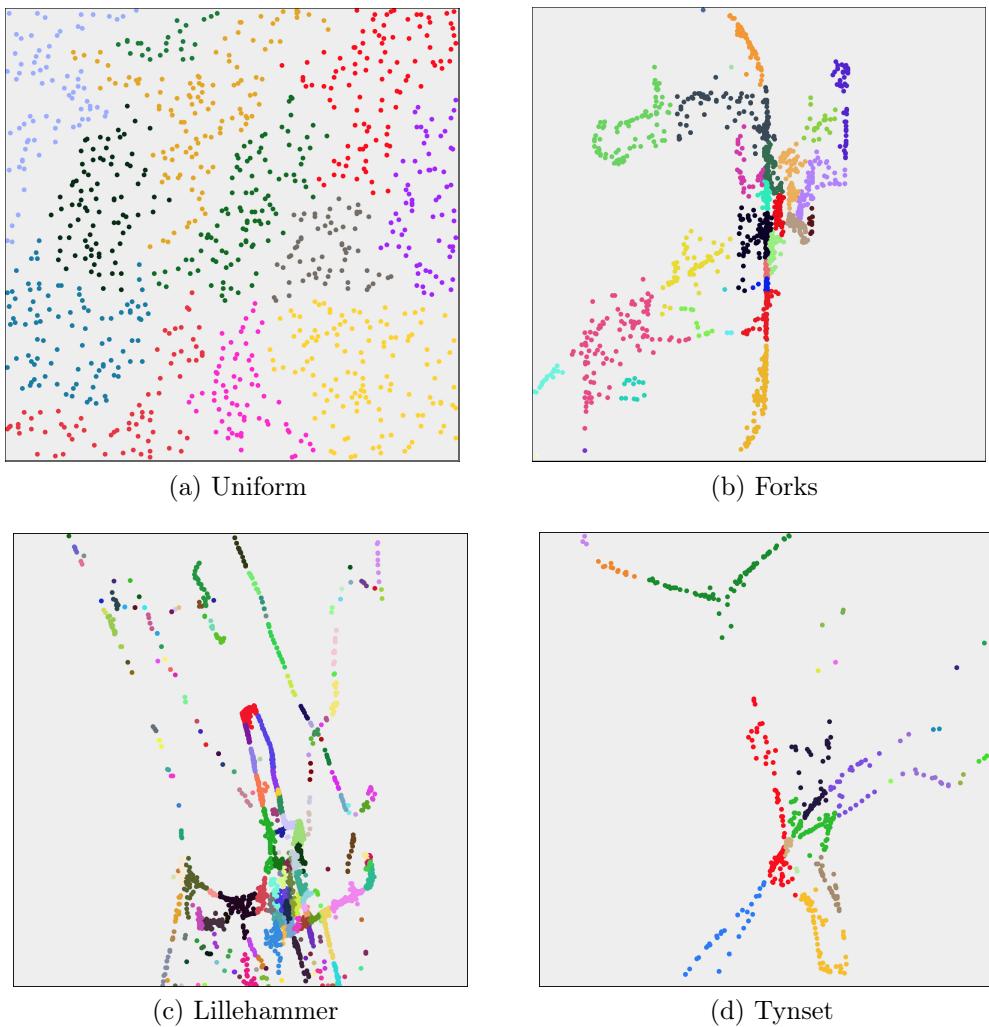


Figure 5.3: K-Nearest Neighbour Clustering on different topologies

5.5.3 Results of on different topologies

Figure 5.3a shows the result of running the K-Nearest Neighbour Clustering algorithm on a uniform distribution topology, while figures 5.3b, 5.3c, 5.3d shows the resulting clusters after running it on topologies on citites and towns collected by WiGLE. The different node colors indicate different group memberships.

5.5.4 Evaluation

By looking at the results of the simulations of the K-Nearest Neighbour Clustering it is obvious that clusters are created. Even in the more chaotic

topologies as Lillehammer, there are distinct clusters that encompasses the closest nodes. However, there are two obvious problems with this algorithm:

- In contrast to agglomerative clustering, our K-Nearest Neighbour Clustering algorithm does not care if the closest distance between the groups is pairwise mutual. This might lead to one group a merging into another group b , even while b has another neighbouring group c that lies closer. If the maximum size K is reached during the merge of a and b , the resulting group will never be able to merge with c even though that would have given a better cluster.
- Once a cluster reaches the maximum size K , there is no possibility for any other nodes to join the cluster. In our simulation of static topologies that might work, but in a real world scenario access points may turn on randomly, and be located in the middle of an already existing cluster. With the current algorithm there is no way to handle this event, and the node could not become a member of the cluster that surrounds it.

In the next section we will look at how we can use group splitting to create a solution for these problems.

5.6 Introducing splitting

To improve the K-Nearest Neighbour Clustering algorithm suggested in the previous section, we can introduce the notion of group splitting. By letting one cluster merge with another without knowing if the merge is mutually beneficial, we are also accepting that not all merges will be optimal. A way to increase the chance of a group ending up with nodes that impacts each other severely would be to accept merges that result in groups larger than K . The group would be too large, but more likely to have the opportunity to merge with the nodes of highest impact. To reduce the group size back to K , the least impactful nodes would have to be kicked out of the group. This is what we call group splitting. This section is dedicated to two different suggestions of how a group could be split.

5.6.1 Introduction to K-means

K-means is a clustering algorithm commonly used for classification purposes in machine learning, usually in the context of unsupervised learning. K-means is an iterative algorithm designed to find a predefined amount of

clusters in a data-set with unclassified data points. The amount of clusters to find is denoted by K (hence K-means).

K centroids are randomly placed somewhere in the data set, then the distance between each data point and each centroid is computed. All data points closest to the same centroid becomes part of the same cluster set. When all data points are associated with a cluster, each cluster computes the mean position of all its data points. The value of the mean position becomes the location of the new centroids. When all the new centroids are computed, one iteration is over. The algorithm continues until the position of all centroids remains unchanged after an iteration. When the position of the centroids are the same two iterations in a row, it means that the solution has converged and all K -clusters has been identified.

A well known problem with K-means is that the result strongly depends on how the algorithm is initiated. This problem can be reduced by running K-means several times and choosing the best result of all the runs.

5.6.2 K-means splitting

K-means clustering is not directly applicable to solve our original clustering challenge. The reason we can not use K-means as the clustering algorithm is because the distributed nature of the nodes and groups. Indeed, a vanilla implementation of K-means requires an extensive overview of the surrounding network topology. Our groups are limited to knowing about their members and their neighbours. It would also be hard to choose a suitable K . Nonetheless, K-means could still potentially help us create better groups. Let us consider the following. We extend the K-Nearest Neighbour Clustering so groups are allowed to merge into a transient group if it exceeds the given maxsize. The purpose of the large transient group is to identify the biggest gap between nodes, and split the group in two new groups with $\text{count}(\text{nodes}) < \text{maxsize}$. If we set $K = 2$, K-means can be used to identify where the transient group should be split to create two more connected clusters.

Randomly picking two nodes to be centroids could lead to undesirable results where the two resulting groups are fundamentally different from the original ones. Recall that what we are trying to do is to reevaluate the division line between the groups and see if K-means can achieve a better split, not create two entirely new clusters. As we already have two independent clusters, we can precalculate their centroids before merging and bootstrap K-means with these centroids instead. If the distance between the most interfering nodes is lower after applying the K-means split, the new groups are applied. On the other hand, if the distance is shorter (meaning a less optimal cut is found) the original groups are restored. As the bootstrapped position of the

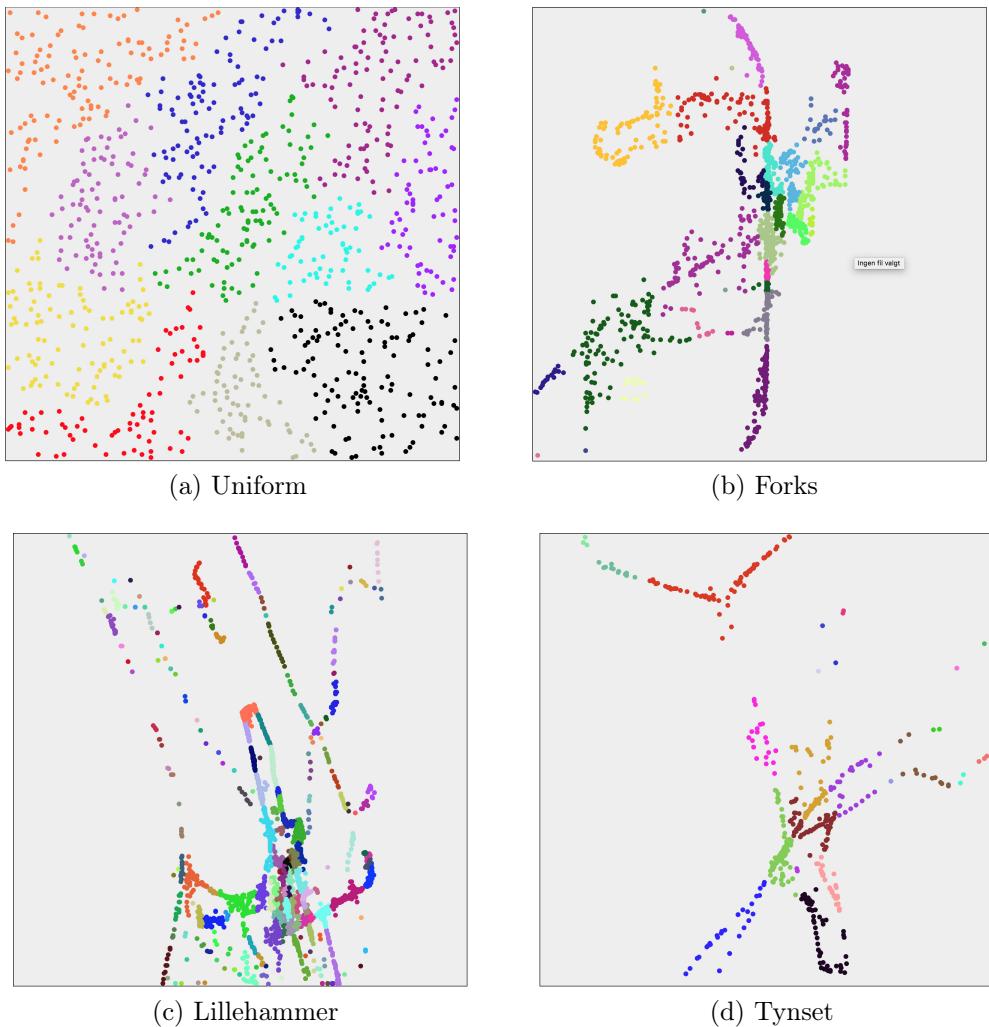


Figure 5.4: K-means splitting on different topologies

centroids is planned and not random, the result is deterministic and we don't have to run the algorithm more than one time to get the result.

5.6.3 Results

The results of simulating the group creation with K-means splitting implemented can be seen in figure 5.4. The groups look very similar to those we computed in 5.3, but in some places there are some major differences. A comparison of the two algorithms executed on the same section of Tynset's town centre can be seen in figure 5.5, and another comparison of a section of the uniform distribution can be seen in figure 5.6.

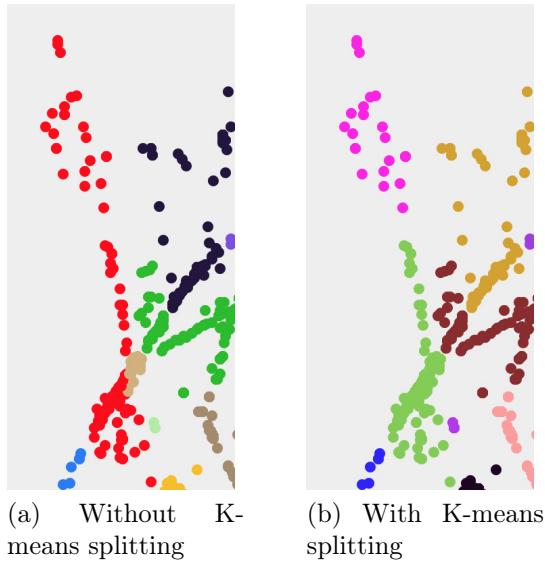


Figure 5.5: Comparison with and without K-means splitting on Tynset

5.6.4 Evaluation

There is little doubt that K-means splitting improves upon the original algorithm, especially when splits only are accepted when the maximum interference of a group is reduced after the split. Even though the clusters in a static environment is slightly improved, the most important use-case for K-means splitting is under the introduction of new nodes to an environment where all groups have converged. Where in the unmodified K-Nearest Neighbour Clustering algorithm there was no specific way to handle a new node in an environment of saturated groups with K-means splitting eventually new groups will be formed on the basis of the updated topology. The weakness of this approach is how ill suited it is to run in a distributed environment. It is trivial to simulate splits using K-means, as during a simulation the absolute coordinates of the nodes in the group are accessible. This is information unavailable for nodes under real circumstances. Hence some major adaptions would have to be made to get this to work in an implementation. One way to realize the algorithm would be to select a physical node as centroid instead of a virtual point. All nodes in sensing range of the surrounding area of the centroid node would know their distance to the centroid. Then, using a relay network to communicate internally in the group, figure out the approximate distance from all nodes not within sensing range of the centroid. It is unknown if such an approach would be accurate enough to work, or have a complexity that is conceivable to implement.

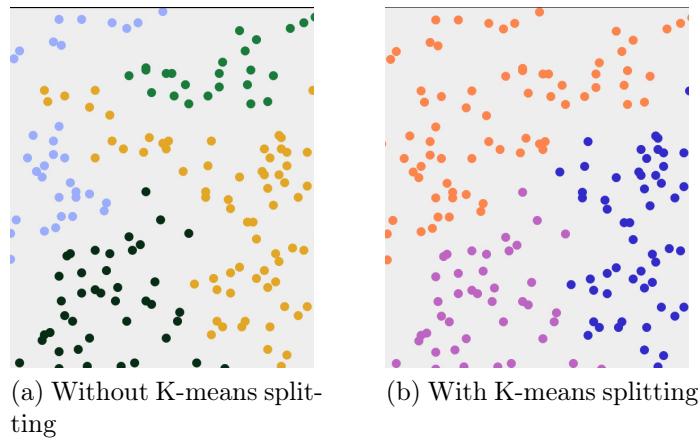


Figure 5.6: Comparison with and without K-means splitting on uniform distribution

5.6.5 Minimum Cut

5.7 Evaluation

In this section we will consider what the results in this chapters can tell us, and cover what questions that still needs answering with respect to distributed group creation and clustering.

5.7.1 Simulation weaknesses

The simulations have taught us about what kind of clustering could work with the assumptions we have set. Even so, there are several questions that still remains, and without considering them we can not expect to get definitive answers of what works or not.

Volatility

All the computations we have done takes a static network topology as an input. The static topology is an image of a network topology in a given state, and is not representative of the volatile nature of modern networks. While it is true that most access points are relatively stationary and constant, it is increasingly popular to create ad-hoc on-demand Wi-Fi networks using cell phones or a computers as access points. Additionally resetting of routers, power outages, etc. means that access points come and go. Hence it is indisputable that network topologies are volatile, and the group clustering

algorithm should, when final, run as a continuous operation. We have considered ways that the algorithm can handle volatility by splitting and continuously create groups, but the simulation framework does not support inserting or removing nodes during the computation, hence only static environments that quickly converges is considered.

Computed interference versus observed interference

Our group computations has been based on the RSSI of all nearby access points, and the dBi values that denotes all RSSIs were calculated using the free space path loss formula. As distance is the only variable that affects the result of our RSSI calculation, it means that the RSSI levels becomes a symmetric binary relation between nodes. In the real world this is not the case. Access points may interfere with each other differently, and the RSSI values may change when there are variations in the environment. This is the reason deadlocks could occur if implementing regular agglomerative hierarchical clustering and the motivation behind creating K-Nearest Neighbour Clustering, even though our topologies would work well with agglomerative clustering. We also have no simulations of how groups would look if the distance between two nodes were not mutual.

Dimensionality

Both the data we have created and the data fetched from `wigle.net` takes the *x* and *y* axes in account. Obtaining data in three dimensions is harder. Wigle places the access points using longitude and latitude as all their data is acquired using ground-level triangulation. By creating clusters based on information in two dimensions we don't know how the algorithm behaves in a 3-dimensional world. It is reasonable to think that an algorithm which is based on distance two dimensions should also work well in three dimensions, but we have done no simulations that can support the claim.

Chapter 6

Group communication and state synchronization

In the previous chapters we have suggested an algorithm for how groups can be formed. In this chapter we will focus on the algorithms and protocols required to enable communication between access points and ultimately groups. There has already been done some work on the subject of access point communication over IP, and we can also rely on previous work in the field of distributed consensus to ensure synchronized group states. We will consider these technologies and suggest a protocol that stitches all required components together.

6.1 Problem overview

In chapter 5 we looked at an algorithm that enables cluster formation in a distributed environment where no node knows the complete layout of the surrounding networks. A number of assumptions were made before we suggested an algorithm. Two of the assumptions encapsulated the three following problems:

- Direct contact between access points is possible
- An underlying group communication protocol is in place
- The state of a group is synchronized throughout all its members

In this chapter we shall attempt to handle these issues to create a more complete scheme for group creation in a distributed environment.

CHAPTER 6. GROUP COMMUNICATION AND STATE SYNCHRONIZATION

6.2 Enabling technologies

When creating the protocol, there is need for some already well-researched technology as a foundation. Raft [17] is a technology created to simplify distributed consensus, and ResFi is a technology created to enable communication over IP between access points. This section will briefly introduce these.

6.2.1 Distributed consensus with Raft

6.2.2 Access point communication with ResFi

6.3 Protocol design

6.4 Protocol implementation

Bibliography

- [1] WiFi Alliance. *WiFi Alliance News Release*. Online: accessed 26-September-2017. 2016. URL: <https://www.wi-fi.org/news-events/newsroom/wi-fi-device-shipments-to-surpass-15-billion-by-end-of-2016>.
- [2] H. Balbi et al. ‘Centralized channel allocation algorithm for IEEE 802.11 networks’. In: *2012 Global Information Infrastructure and Networking Symposium (GIIS)*. Dec. 2012, pp. 1–7. DOI: 10.1109/GIIS.2012.6466657.
- [3] G. Bianchi. ‘Performance analysis of the IEEE 802.11 distributed co-ordination function’. In: *IEEE Journal on Selected Areas in Communications* 18.3 (Mar. 2000), pp. 535–547. ISSN: 0733-8716. DOI: 10.1109/49.840210.
- [4] Robert G. Chamberlain. *Computing Distances*. Online: accessed 14-Nov-2017. 1999. URL: <https://cs.nyu.edu/visual/home/proj/tiger/gisfaq.html>.
- [5] Cisco. *Radio Resource Management under Unified Wireless Networks*. Accessed 16th Nov, 2017. May 2010. URL: <https://www.cisco.com/c/en/us/support/docs/wireless-mobility/wireless-lan-wlan/71113-rrm-new.html>.
- [6] 2017 Cisco VNI. ‘Cisco Visual Networking Index: Forecast and Methodology, 2016–2021’. In: (June 2017). URL: <https://www.cisco.com/c/en/us/solutions/collateral/service-provider/visual-networking-index-vni/complete-white-paper-c11-481360.html>.
- [7] Python Software Foundation. *Python 3*. Online: accessed 05-July-2017. 2017. URL: <http://python.org>.
- [8] H. T. Friis. ‘A Note on a Simple Transmission Formula’. In: *Proceedings of the Institute of Radio Engineers* 34.5 (May 1946), pp. 254–256. ISSN: 0096-8390. DOI: 10.1109/JRPROC.1946.234568. URL: <http://dx.doi.org/10.1109/JRPROC.1946.234568>.

BIBLIOGRAPHY

- [9] Gartner. *Gartner Press Release*. Online: accessed 26-September-2017. 2017. URL: <http://www.gartner.com/newsroom/id/3598917>.
- [10] Ramakrishna Gummadi et al. ‘Understanding and Mitigating the Impact of RF Interference on 802.11 Networks’. In: *SIGCOMM Comput. Commun. Rev.* 37.4 (Aug. 2007), pp. 385–396. ISSN: 0146-4833. DOI: 10.1145/1282427.1282424. URL: <http://doi.acm.org/10.1145/1282427.1282424>.
- [11] ‘IEEE Standard for Information Technology- Telecommunications and Information Exchange Between Systems- Local and Metropolitan Area Networks- Specific Requirements Part 11: Wireless LAN Medium Access Control (MAC) and Physical Layer (PHY) Specifications Amendment 7: 4.9 GHz-5 GHz Operation in Japan’. In: *IEEE Std 802.11j-2004* (2004), pp. 1–40. DOI: 10.1109/IEEESTD.2004.95388.
- [12] *The JSON Data Interchange Format*. Tech. rep. Standard ECMA-404 1st Edition / October 2013. ECMA, Oct. 2013. URL: <http://www.ecma-international.org/publications/files/ECMA-ST/ECMA-404.pdf>.
- [13] D. Müllner. ‘Modern hierarchical, agglomerative clustering algorithms’. In: *ArXiv e-prints* (Sept. 2011). arXiv: 1109.2378 [stat.ML].
- [14] Rohan Murty et al. ‘Designing High Performance Enterprise Wi-Fi Networks’. In: *Proceedings of the 5th USENIX Symposium on Networked Systems Design and Implementation*. NSDI’08. San Francisco, California: USENIX Association, 2008, pp. 73–88. ISBN: 111-999-5555-22-1. URL: <http://dl.acm.org/citation.cfm?id=1387589.1387595>.
- [15] Rohan Murty et al. ‘Dyson: An Architecture for Extensible Wireless LANs’. In: *Proceedings of the 2010 USENIX Conference on USENIX Annual Technical Conference*. USENIXATC’10. Boston, MA: USENIX Association, 2010, pp. 15–15. URL: <http://dl.acm.org/citation.cfm?id=1855840.1855855>.
- [16] Aerohive Networks. ‘Radio Resource Management in HiveOS’. In: (2011). Technical report. URL: http://docs.aerohive.com/pdfs/Aerohive-Solution_Brief-RadioResource_Management_in_HiveOS.pdf.
- [17] Diego Ongaro and John Ousterhout. ‘In Search of an Understandable Consensus Algorithm’. In: *Proceedings of the 2014 USENIX Conference on USENIX Annual Technical Conference*. USENIX ATC’14. Philadelphia, PA: USENIX Association, 2014, pp. 305–320. ISBN: 978-1-

BIBLIOGRAPHY

- 931971-10-2. URL: <http://dl.acm.org/citation.cfm?id=2643634.2643666>.
- [18] Roger W Sinnott. ‘Virtues of the Haversine’. In: *Sky and Telescope* 68.2 (1984), p. 158.
- [19] Lalith Suresh et al. ‘Towards Programmable Enterprise WLANS with Odin’. In: *Proceedings of the First Workshop on Hot Topics in Software Defined Networks*. HotSDN ’12. Helsinki, Finland: ACM, 2012, pp. 115–120. ISBN: 978-1-4503-1477-0. DOI: 10.1145/2342441.2342465. URL: <http://doi.acm.org/10.1145/2342441.2342465>.
- [20] Sven Zehl (szehl). *ResFi*. <https://github.com/resfi/resfi>. 2016.
- [21] WiGLE. *WiGLE API*. Nov. 2017. URL: <https://api.wigle.net/swagger>.
- [22] WiGLE. *WiGLE: Wireless Network Mapping*. July 2017. URL: <https://wigle.net/>.
- [23] Wikipedia. *Free-space path loss — Wikipedia, The Free Encyclopedia*. Accessed 4th July, 2017. 2017. URL: <http://en.wikipedia.org/w/index.php?title=Free-space%5C%20path%5C%20loss&oldid=765763753>.
- [24] Sven Zehl et al. ‘ResFi: A Secure Framework for Self Organized Radio Resource Management in Residential WiFi Networks’. In: *17th IEEE International Symposium on a World of Wireless, Mobile and Multimedia Networks (IEEE WoWMoM 2016)*. accepted for publication. Coimbra, Portugal, June 2016. URL: <http://www.tkn.tu-berlin.de/fileadmin/fg112/Papers/2016/zehl16resfi.pdf>.